

Relational data

- SQLite
- pandas



Two tables

Table: country			
name	population	area	capital
'Denmark'	5748769	42931	'Copenhagen'
'Germany'	82800000	357168	'Berlin'
'USA'	325719178	9833520	'Washington, D.C.'
'Iceland'	334252	102775	'Reykjavik'

Table: city			
name	country	population	established
'Copenhagen'	'Denmark'	775033	800
'Aarhus'	'Denmark'	273077	750
'Berlin'	'Germany'	3711930	1237
'Munich'	'Germany'	1464301	1158
'Reykjavik'	'Iceland'	126100	874
'Washington D.C.'	'USA'	693972	1790
'New Orleans'	'USA'	343829	1718
'San Francisco'	'USA'	884363	1776

SQL

pronounced ,ɛs,kjuː'ɛl or 'siːkwəl

- SQL = Structured Query **Language**
- **Database = collection of tables** stored persistently on disk
- ANSI and ISO standards since 1986 and 1987, respectively; origin early 70s
- Widespread used SQL databases (can handle many tables/rows/users):
[Oracle](#), [MySQL](#), [Microsoft SQL Server](#), [PostgreSQL](#) and [IBM DB2](#)
- **SQLite** is a very lightweight version storing a database in a **single file**, without a separate database server; first release in 2000
- SQLite is included in both iOS and Android mobile phones and operating systems like Mac OS, Windows, and Fedora Linux

Table: country			
name	population	area	capital
'Denmark'	5748769	42931	'Copenhagen'
'Germany'	82800000	357168	'Berlin'
'USA'	325719178	9833520	'Washington, D.C.'
'Iceland'	334252	102775	'Reykjavik'



The Course "[Databases](#)" gives a more in-depth introduction to SQL (MySQL)

Examples : SQL statements

Table: country			
name	population	area	capital
'Denmark'	5748769	42931	'Copenhagen'
'Germany'	82800000	357168	'Berlin'
'USA'	325719178	9833520	'Washington, D.C.'
'Iceland'	334252	102775	'Reykjavik'

- CREATE TABLE country (name, population, area, capital)
- INSERT INTO country VALUES ('Denmark', 5748769, 42931, 'Copenhagen')
- UPDATE country SET population=5748770 WHERE name='Denmark'
- SELECT name, capital FROM country WHERE population >= 1000000
> [('Denmark', 'Copenhagen'), ('Germany', 'Berlin'), ('USA', 'Washington, D.C.')]
- SELECT * FROM country WHERE capital = 'Berlin'
> [('Germany', 82800000, 357168, 'Berlin')]
- SELECT country.name, city.name, city.established FROM city, country
WHERE city.name=country.capital AND city.population < 700000
> [('Iceland', 'Reykjavik', 874), ('USA', 'Washington, D.C.', 1790)]
- DELETE FROM country WHERE name = 'Germany'
- DROP TABLE country

sqlite-example.py

```
import sqlite3

connection = sqlite3.connect('example.sqlite') # creates file if necessary
c = connection.cursor()

c.executescript('''DROP TABLE IF EXISTS country; -- multiple SQL statements
                  DROP TABLE IF EXISTS city''')

countries = [('Denmark', 5748769, 42931, 'Copenhagen'),
             ('Germany', 82800000, 357168, 'Berlin'),
             ('USA', 325719178, 9833520, 'Washington, D.C.'),
             ('Iceland', 334252, 102775, 'Reykjavik')]

cities = [('Copenhagen', 'Denmark', 775033, 800),
          ('Aarhus', 'Denmark', 273077, 750),
          ('Berlin', 'Germany', 3711930, 1237),
          ('Munich', 'Germany', 1464301, 1158),
          ('Reykjavik', 'Iceland', 126100, 874),
          ('Washington, D.C.', 'USA', 693972, 1790),
          ('New Orleans', 'USA', 343829, 1718),
          ('San Francisco', 'USA', 884363, 1776)]

c.execute('CREATE TABLE country (name, population, area, capital)')
c.execute('CREATE TABLE city (name, country, population, established)')
c.executemany('INSERT INTO country VALUES (?, ?, ?, ?)', countries)
c.executemany('INSERT INTO city VALUES (?, ?, ?, ?)', cities)

connection.commit() # save data to database before closing
connection.close()
```

SQLite

SQLite query examples

try to avoid using the asterisk (*) as a good habit

www.sqlitetutorial.net/sqlite-select

sqlite-example.py (continued)

```
for row in c.execute('SELECT * FROM country'): # * = all columns, execute returns iterator
    print(row)                                # row is by default a Python tuple

for row in c.execute('''SELECT * FROM city, country -- all pairs of rows from city x country
                        WHERE city.name = country.capital AND city.population < 700000'''):
    print(row)

print(*c.execute('''SELECT country.name,
                    COUNT(city.name) AS cities,
                    100 * SUM(city.population) / country.population
                    FROM city JOIN country ON city.country = country.name -- SQL join 2 tables
                    WHERE city.population > 500000 -- only consider big cities
                    GROUP BY city.country -- output has one row per group of rows
                    ORDER BY cities DESC, SUM(city.population) DESC''')) # ordering of output
```

Python shell

```
| ('Denmark', 5748769, 42931, 'Copenhagen')
| ('Germany', 82800000, 357168, 'Berlin')
| ('USA', 325719178, 9833520, 'Washington, D.C.')
| ('Iceland', 334252, 102775, 'Reykjavik')

| ('Reykjavik', 'Iceland', 126100, 874, 'Iceland', 334252, 102775, 'Reykjavik')
| ('Washington, D.C.', 'USA', 693972, 1790, 'USA', 325719178, 9833520, 'Washington, D.C.')
| ('Germany', 2, 6) ('USA', 2, 0) ('Denmark', 1, 13)
```

SQL injection

Right way

```
c.execute('INSERT INTO users VALUES (?)', (user,))
```

unsafe-example.py

```
import sqlite3
connection = sqlite3.connect('users.sqlite')
c = connection.cursor()
c.execute('CREATE TABLE users (name)')
while True:
    user = input('New user: ')
    c.executescript(f'INSERT INTO users VALUES ("{user}")')
    connection.commit()
    print(list(c.execute('SELECT * FROM users')))
```

can execute a string
containing several
SQL statements



Insecure
NEVER use f-string
with user input

Python shell

```
> New user: gerth
| [('gerth',)]
> New user: guido
| [('gerth',), ('guido',)]
> New user: evil"); DROP TABLE users; --
| sqlite3.OperationalError: no such table: users
```

INSERT INTO users VALUES ("evil"); DROP TABLE users; --")

HI, THIS IS
YOUR SON'S SCHOOL.
WE'RE HAVING SOME
COMPUTER TROUBLE.



OH, DEAR - DID HE
BREAK SOMETHING?
IN A WAY -



DID YOU REALLY
NAME YOUR SON
Robert'); DROP
TABLE Students;-- ?



OH. YES. LITTLE
BOBBY TABLES,
WE CALL HIM.

WELL, WE'VE LOST THIS
YEAR'S STUDENT RECORDS.
I HOPE YOU'RE HAPPY.




AND I HOPE
YOU'VE LEARNED
TO SANITIZE YOUR
DATABASE INPUTS.

Pandas



- Comprehensive Python library for data manipulation and analysis, in particular tables and time series
- Pandas **data frames** = tables
- Supports interaction with SQL, CSV, JSON, ...
- Integrates with Jupyter, numpy, matplotlib, ...

pandas
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$



Pandas integration with Jupyter

- Tables (Pandas data frames) are rendered nicely in Jupyter

```
In [1]: 1 import pandas as pd
        2 students = pd.read_csv('students.csv')
        3 students
```

Out [1]:

	Name	City
0	Donald Duck	Copenhagen
1	Goofy	Aarhus
2	Mickey Mouse	Aarhus

students.csv

```
Name, City
"Donald Duck", "Copenhagen"
"Goofy", "Aarhus"
"Mickey Mouse", "Aarhus"
```

Reading tables (data frames)

- Pandas provides functions for reading and writing pandas.DataFrames as files in different data formats, e.g. SQLite, csv and EXCEL files

pandas-example.py

```
import pandas as pd
import sqlite3
connection = sqlite3.connect('example.sqlite')
countries = pd.read_sql_query('SELECT * FROM country', connection)
cities = pd.read_sql_query('SELECT * FROM city', connection)
students = pd.read_csv('students.csv')
students.to_sql('students', connection, if_exists='replace')
students.to_excel('students.xlsx', sheet_name='students')
print(students)
```

Python shell

	Name	City
0	Donald Duck	Copenhagen
1	Goofy	Aarhus
2	Mickey Mouse	Aarhus

	A	B	C	D
1		Name	City	
2	0	Donald Duck	Copenhagen	
3	1	Goofy	Aarhus	
4	2	Mickey Mouse	Aarhus	
5				

Selecting columns and rows

Table: country			
name	population	area	capital
'Denmark'	5748769	42931	'Copenhagen'
'Germany'	82800000	357168	'Berlin'
'USA'	325719178	9833520	'Washington, D.C.'
'Iceland'	334252	102775	'Reykjavik'

Python shell

```
> countries['name']           # select column
> countries.name              # same as above
> countries[['name', 'capital']] # select multiple columns, note double-[]
> countries.head(2)           # first 2 rows
> countries[1:3]               # slicing rows, rows 1 and 2
> countries[::2]               # slicing rows, rows 0 and 2
> countries.at[1, 'area']      # indexing cell by (row label, column name)
> cities[(cities['name'] == 'Berlin') | (cities['name'] == 'Munich')] # select rows
|      name  country  population  established
|  2  Berlin  Germany    3711930         1237 # note original row labels
|  3  Munich  Germany    1464301         1158
> pd.DataFrame([[1, 2], [3, 4], [5, 6]], columns=['x', 'y']) # create DF from list
> pd.DataFrame(np.random.random((3, 2)), columns=['x', 'y']) # from numpy
```

Row labels

Python shell

```
> df = pd.DataFrame(np.arange(1, 13).reshape(3, 4),
                    index=['q', 'w', 'e'],          # row labels
                    columns=['c', 'a', 'd', 'e'])    # column names

> df
|   c   a   d   e
| q   1   2   3   4 # row labels can be strings
| w   5   6   7   8
| e   9  10  11  12

> df.loc['w':'e', ['e', 'a']] # slice of labeled rows
|      e   a
| w     8   6
| e    12  10

> df.loc['w'] # single row (a one-dimensional pd.Series)
| c     5
| a     6
| d     7
| e     8
| Name: w, dtype: int32

> df.iloc[:2, :2] # use iloc to work with integer indexes
|   c   a
| q   1   2
| w   5   6
```

Merging tables and creating a new column

pandas-example.py


```
M = pd.merge(countries, cities, left_on='capital', right_on='name')
# both data frames had a 'name' and 'population' column
M1 = M.rename(columns={
    'population_x': 'country_population',
    'population_y': 'capital_population'
})
M2 = M1.drop(columns=['name_x', 'name_y'])
M2['%pop in capital'] = M2.capital_population / M2.country_population # add column
M2.sort_values('%pop in capital', ascending=False, inplace=True)
print(M2[['country', '%pop in capital']])
```

Python shell

	country	%pop in capital	
3	Iceland	0.377260	# note row labels are permuted
0	Denmark	0.134817	
1	Germany	0.044830	
2	USA	0.002131	

Pandas and Matplotlib

- Dataframes have a `.plot` method (using `matplotlib.pyplot`)
- Module `yfinance` provides access to Yahoo! Finance's API as Pandas dataframes
- `pip install yfinance`



pandas-yfinance.py

```
import matplotlib.pyplot as plt
import yfinance
df = yfinance.download(['AAPL', 'GOOGL', 'MSFT', 'ZM'], start='2000-01-01', interval='1d')
df['Close'].plot()
plt.legend()
plt.show()
```

Hierarchical / Multi-level indexing (MultiIndex)

Python shell

```
> df.tail(2)
| Price          Close          ...          Volume
| Ticker          AAPL          GOOGL          MSFT          ...          GOOGL          MSFT          ZM
| Date
| 2026-04-23    273.429993    338.890015    415.750000    ...    18650300.0    38308000    3810800.0
| 2026-04-24    271.059998    344.399994    424.619995    ...    26400000.0    27413900    4252300.0
> df['Close'].tail(2)
| Ticker          AAPL          GOOGL          MSFT          ZM
| Date
| 2026-04-23    273.429993    338.890015    415.750000    90.010002
| 2026-04-24    271.059998    344.399994    424.619995    92.029999
> df['Close']['GOOGL'].tail(2)
| Date
| 2026-04-23          338.890015
| 2026-04-24          344.399994
| Name: GOOGL, dtype: float64
> df.loc[:, pd.IndexSlice[:, 'GOOGL']].tail(2)
| Price          Close          High          Low          Open          Volume
| Ticker          GOOGL          GOOGL          GOOGL          GOOGL          GOOGL
| Date
| 2026-04-23    338.890015    341.959991    336.179993    341.179993    18650300.0
| 2026-04-24    344.399994    345.269989    335.390015    338.730011    26400000.0
```

Both rows and columns
can have multi-level
indexing

Python shell

```
> df.columns
| MultiIndex([( 'Close',  'AAPL'),
|             ( 'Close',  'GOOGL'),
|             ( 'Close',  'MSFT'),
|             ( 'Close',   'ZM'),
|             ( 'High',   'AAPL'),
|             ( 'High',  'GOOGL'),
|             ( 'High',  'MSFT'),
|             ( 'High',   'ZM'),
|             ( 'Low',   'AAPL'),
|             ( 'Low',  'GOOGL'),
|             ( 'Low',  'MSFT'),
|             ( 'Low',   'ZM'),
|             ( 'Open',  'AAPL'),
|             ( 'Open',  'GOOGL'),
|             ( 'Open',  'MSFT'),
|             ( 'Open',   'ZM'),
|             ( 'Volume', 'AAPL'),
|             ( 'Volume', 'GOOGL'),
|             ( 'Volume', 'MSFT'),
|             ( 'Volume',  'ZM')]),
|          names=['Price', 'Ticker'])
```